


Assignment 1 – Part 2: Logistic Regression and Neural Networks

1. Logistic Regression


The formula for the given hypothesis in logistic regression is a logistic/sigmoid given below:

```
output = 1 / (1 + np.exp(-z))  
return output
```

After plotting the plot_sigmoid.py function to plot the sigmoid function is given below:



Note: After running the `plot_data.py` we can clearly observe that the data have been normalised, to enable the data to be more easily optimised.



In `plot_boundary` I am rearranging the terms in the equation of the hypothesis function.

```
def plot_boundary(X, theta, ax1):  
  
    a = X[..., 1]  
    b = X[..., 2]  
  
    min_x1 = np.max(a)  
    max_x1 = np.min(a)  
    ia = np.where(a == min_x1)  
    ib = np.where(a == max_x1)  
    x2_on_min_x1 = b[ia]  
    x2_on_max_x1 = b[ib]
```

```
x_array = np.array([min_x1, max_x1])
y_array = np.array([x2_on_min_x1,
x2_on_max_x1])
ax1.plot(x_array, y_array, c='black',
label='decision boundary')
```

plots are given below after running the ml_assgn1_ex1 for

```
alpha = 0.024
iterations = 100
```



```
cost = (-output)*np.log(hypothesis) - (1 -
output) * np.log(1-hypothesis)
```


For training sample 20

Final training cost: 0.17416

Minimum training cost: 0.17416, on iteration #100

Final test cost: 0.73745








For training sample 40

Final training cost: 0.41818

Minimum training cost: 0.41818, on iteration #100

Final test cost: 0.39502








For sample 65

Final training cost: 0.42026

Minimum training cost: 0.42026, on iteration #60

Final test cost: 0.02577








For sample 77

Final training cost: 0.42026

Minimum training cost: 0.42026, on iteration #60

Final test cost: 0.02577






The training set must be separate from the test set whereas the training phase consumes the training set. In order to find a set of parameter values that minimise a certain cost function over the whole training set. In the case of test set it is for testing the model to check how it performs on new unseen versions of the same classes.

```
x1 = X[:, 0]
x2 = X[:, 1]
X = np.column_stack((X, x1*x2))
X = np.column_stack((X, x1**2))
X = np.column_stack((X, x2**2))
```

```
alpha = 0.01

iterations = 100
```

```
theta = np.zeros((6))
```



In task 8 I am modifying the function `gradient_descent_training.py` to store the current cost for training set and testing set.

I am storing the cost of the training set to `cost_vector_train` and for the test set to `cost_vector_test`.

```
for j in range(len(theta))
```

```
hypothesis = calculate_hypothesis(X,  
theta_temp, i)
```

```
#####/  
sigma[j] += (hypothesis - output) * X[i, j]
```

Here I have updated the temp

```
theta_temp[j] = theta_temp[j] - alpha * sigma[j]
```

after running the ml_assgn_4.py file I can see the following observations:


sample 20

Final train cost: 0.43985

Minimum train cost: 0.43985, on iteration #100

Final test cost: 0.40566

Minimum test cost: 0.40566, on iteration #100




sample 30

Final train cost: 0.35294

Minimum train cost: 0.35294, on iteration #100

Final test cost: 0.49431

Minimum test cost: 0.45083, on iteration #16




sample 40

Final train cost: 0.26615

Minimum train cost: 0.26615, on iteration #100

Final test cost: 0.65850

Minimum test cost: 0.51458, on iteration #7




sample 70

Final train cost: 0.40522

Minimum train cost: 0.40522, on iteration #100

Final test cost: 0.40428

Minimum test cost: 0.39321, on iteration #7



In assgn1_ex5.py I am adding extra features and analysing the effect.

```
def gradient_descent_training(X_train, y_train,  
X_test, y_test, theta, alpha, iterations):
```

```
    output1 = y_train[i]  
    hypothesis = calculate_hypothesis(X_train,  
    theta_temp, i)
```




```
for j in range(len(theta)):
```

```
    output1 = y_train[i]  
    hypothesis = calculate_hypothesis(X_train,  
    theta_temp, i)
```


```
    sigma[j] += (hypothesis - output1) * X_train[i,  
    j]
```

```
it_cost_vector_train = compute_cost(X_train,  
y_train, theta)  
it_cost_vector_test = compute_cost(X_test,  
y_test, theta)  
  
cost_vector_train =  
np.append(cost_vector_train,  
it_cost_vector_train)  
cost_vector_test = np.append(cost_vector_test,  
it_cost_vector_test)
```

Final train cost: 0.36887
Minimum train cost: 0.36887, on iteration #100
Final test cost: 0.42051
Minimum test cost: 0.41684, on iteration #37




Sample 40




Final train cost: 0.42920
Minimum train cost: 0.42920, on iteration #100
Final test cost: 0.41132
Minimum test cost: 0.40567, on iteration #19

Sample 60




Final train cost: 0.39530
Minimum train cost: 0.39530, on iteration #100
Final test cost: 0.44090
Minimum test cost: 0.42626, on iteration #10

Sample 70



Final train cost: 0.43500
Minimum train cost: 0.43500, on iteration #100
Final test cost: 0.17417
Minimum test cost: 0.17385, on iteration #56

The logistic regression unit cannot solve the XOR classification problem because the classes in XOR are not linearly separable.



2. Neural Network

In my first case I am modifying sigmoid.py to use the sigmoid function.

```
output = 1.0 / (1.0 + np.exp(-z))  
return output
```

```
def backward_pass(self, inputs, targets,  
learning_rate):
```

Step 1: In step 1 the output deltas are used to update the weights of the output layer.

```
output_deltas = (outputs - targets) *  
sigmoid_derivative(sigmoid(outputs[i]))
```

Step 2: Hidden deltas are used to update the weights of the hidden layer.

Note: Here I need to backpropagate the error to the hidden neurons

```
hidden_deltas[i] =  
sigmoid_derivative(sigmoid(self.y_hidden[i]))*sigma
```

Step 3: Her I am updating the output weights, the connections from the hidden neurons to the output neurons.


```
self.w_out[i,j] = self.w_out[i,j]-  
learning_rate*output_deltas[j]*sigmoid(self.y_hidden[  
i])
```

Step4: Here I am updating the hidden weights, the connections from the hidden neurons to the inputs.


```
self.w_hidden[i, j] = self.w_hidden[i, j] -  
learning_rate * hidden_deltas[j] * sigmoid(inputs[i])
```

Learning rate	Cost Function
---------------	---------------


0.1




0.05



1.0



3



In task 2 I am checking for the NOR gate and with the truth table given below:

2 Input NOR gate		
A	B	$\overline{A+B}$
0	0	1
0	1	0
1	0	0
1	1	0

```
y = np.array([1, 0, 0, 0])
```

For the NOR gate output values I am getting the following samples of target and predicted value.


Sample #01 | Target value: 1.00 | Predicted value: 0.97625

Sample #02 | Target value: 0.00 | Predicted value: 0.00889

Sample #03 | Target value: 0.00 | Predicted value: 0.00905

Sample #04 | Target value: 0.00 | Predicted value: 0.00002

Minimum cost: 0.00037, on iteration #10000




Now for AND gate with the truth table given below:

2 Input AND gate		
A	B	A.B
0	0	0
0	1	0
1	0	0
1	1	1

```
y = np.array([0, 0, 0, 1])
```

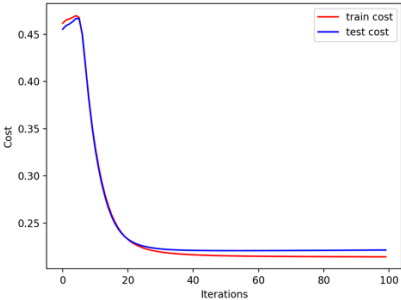
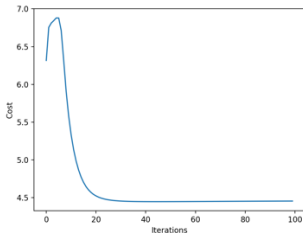
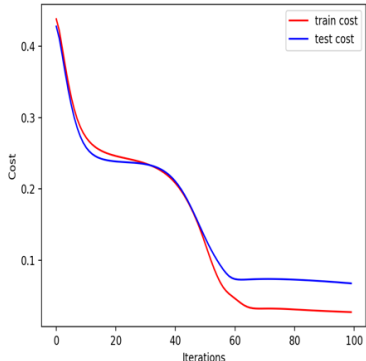
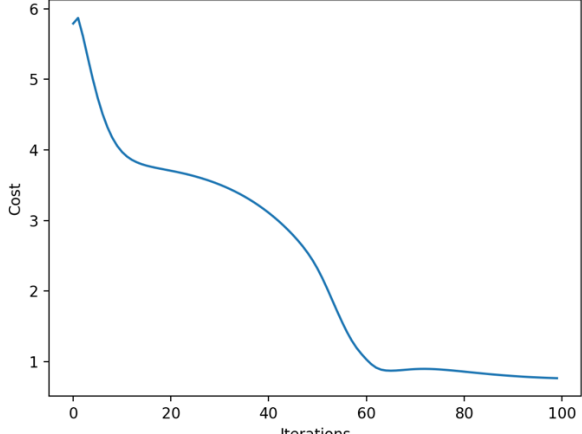
For the AND gate output values I am getting the following samples of target and predicted value.

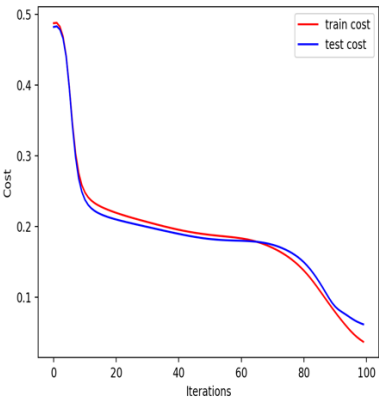
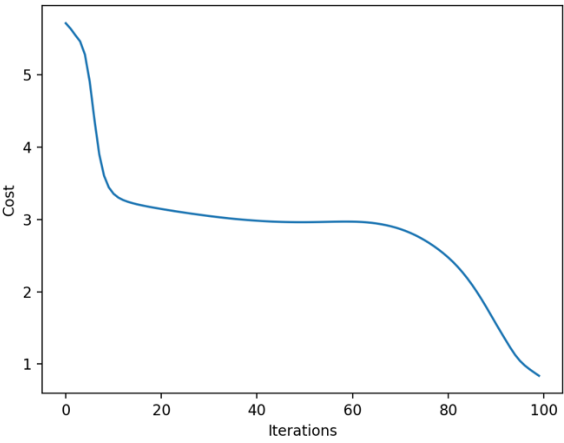
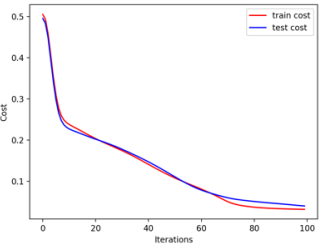
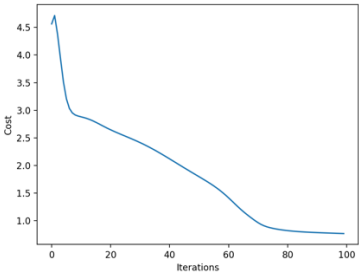
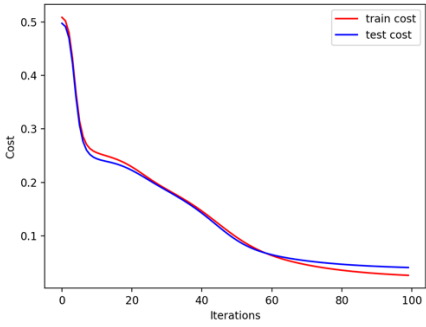
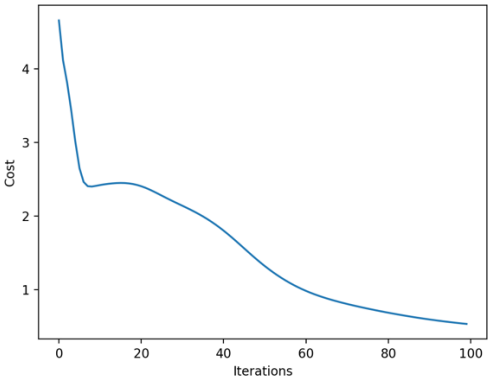
Sample #01 | Target value: 0.00 | Predicted value: 0.00032
Sample #02 | Target value: 0.00 | Predicted value: 0.01463
Sample #03 | Target value: 0.00 | Predicted value: 0.01463
Sample #04 | Target value: 1.00 | Predicted value: 0.96725
Minimum cost: 0.00085, on iteration #10000

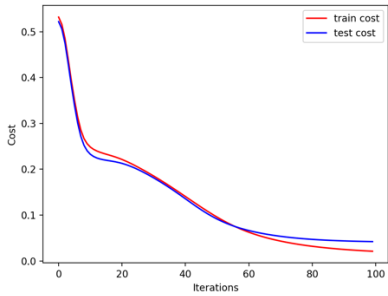
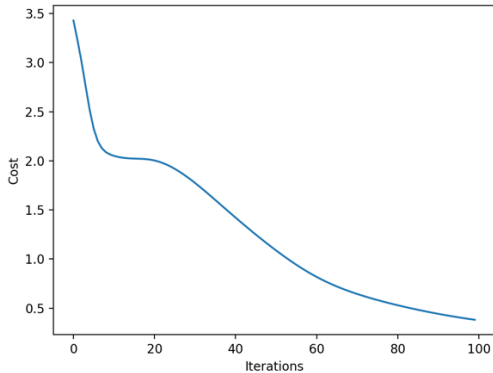


NOTE

- One versus all classification is a neat recipe to build a multiway classifier using just a series of binary classifiers.
- Use the one versus rest classification option makes it challenging to handle large datasets due to many class instances.

Hidden Neurons	Train Cost Test Code	Cost function
1		
2		
3		

	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (0.1 to 0.5). The legend indicates 'train cost' (red line) and 'test cost' (blue line). Both costs decrease sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (1 to 5). The cost decreases sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>
5	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (0.1 to 0.5). The legend indicates 'train cost' (red line) and 'test cost' (blue line). Both costs decrease sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (1.0 to 4.5). The cost decreases sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>
7	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (0.1 to 0.5). The legend indicates 'train cost' (red line) and 'test cost' (blue line). Both costs decrease sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>	 <p>Plot of Cost vs Iterations. The x-axis is labeled 'Iterations' (0 to 100) and the y-axis is labeled 'Cost' (1 to 4). The cost decreases sharply from iteration 0 to 20, then more gradually, reaching a minimum around iteration 100.</p>

10	 <p>A line graph with 'Iterations' on the x-axis (0 to 100) and 'Cost' on the y-axis (0.0 to 0.5). Two lines are plotted: 'train cost' (red) and 'test cost' (blue). Both lines start at approximately 0.55 at iteration 0. The train cost decreases steadily to about 0.05 at iteration 100. The test cost follows a similar path but remains slightly higher than the train cost after iteration 20, ending at approximately 0.08 at iteration 100.</p>	 <p>A line graph with 'Iterations' on the x-axis (0 to 100) and 'Cost' on the y-axis (0.5 to 3.5). A single blue line is plotted, starting at approximately 3.4 at iteration 0. It drops sharply to about 2.1 by iteration 10, then continues to decrease more gradually, reaching approximately 0.4 at iteration 100.</p>

For hidden neuron 10 is the best for iteration.

We haven't generalized so we don't need these many basic neurons and it is fitting my training set well but it is not generalized.